

Numerical Optimization for ML&DL (NOFML&DL)

Artificial Intelligence and Data Science

شرح بالعربي

1

Agenda

GD Applied to Multivariate LR.

Features Scaling.

x	y
Area	Price
8450	208500
9600	181500
11250	223500
9550	140000
14260	250000
14115	143000
10084	307000

x_0	x_1	x_2	x_3	x_4	y
Bias (intersect) multiplier	Income	House Age	Number of Rooms	Number of Bedrooms	Price (e+06)
1	79545	5	7	4	1.059
1	79248	6	6	3	1.505
1	61287	5	8	5	1.058
1	63345	7	5	3	1.260
1	59982	5	7	4	6.309

GD Applied to Multivariate LR

- Single Variable LR: $h_{\theta}(x) = \theta_0 + \theta_1 x$
- Multi Variable LR
- $h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$
- $x_0 = 1$
- $h_{\theta}(x) = \Theta^T X$

Hypothesis: $h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$

Parameters: $\theta_0, \theta_1, \dots, \theta_n$

Cost function:

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Gradient descent:

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n)$$

} (simultaneously update for every $j = 0, \dots, n$)

GD Applied to Multivariate LR

Gradient Descent

Previously (n=1):

Repeat {

$$\theta_0 := \theta_0 - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})}_{\frac{\partial}{\partial \theta_0} J(\theta)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

(simultaneously update θ_0, θ_1)

}

New algorithm ($n \geq 1$):

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

(simultaneously update θ_j for
 $j = 0, \dots, n$)

}

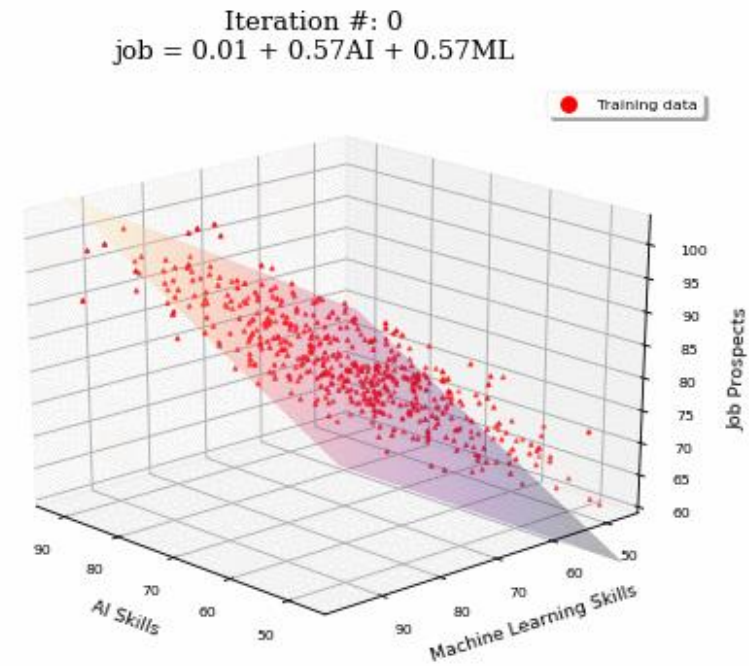
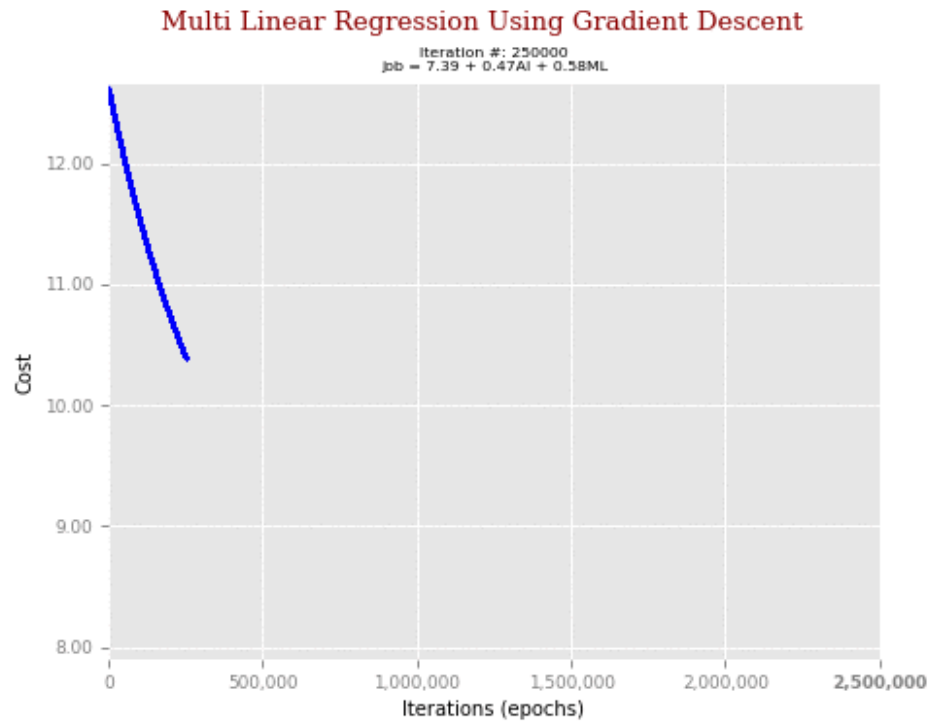
$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$$

$$\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_2^{(i)}$$

...


GD Applied
to
Multivariate
LR



GD Applied to Multivariate LR

Features Scaling:


- **Features Scaling:**
 - Make sure features are on similar scale.
 - For gradient-based algorithms, features scaling improves the convergence speed.



Features Scaling:


- **Problem Statement:**

- If we have two features x_1 with large scale and x_2 with low scale the equivalent θ_1 will be small and have a small search range and θ_2 will be large and has large search range.
- In the opposite side, as long as the gradient vector components depends on the feature value, the gradient component for the large-scale feature will be larger than the small.




Features Scaling:

- **Features Statement:**
 - This implies a small range of θ_1 with large update in its direction and large range of θ_2 with small update in its direction.
 - This makes the gradient descent oscillates during training and consumes large number of iterations.




Features Scaling:

- Make sure features are on similar scale.
- For gradient-based algorithms, features scaling improves the convergence speed.
- Distance-based algorithms like **KNN**, **K-means**, and **SVM** are most affected by the range of features.
- **Tree-based** algorithms, on the other hand, are fairly **insensitive** to the scale of the features.



Features Scaling:

- Use feature scaling when the algorithm calculates distances (K-Nearest Neighbor and Support Vector Machines) or is trained with Gradient Descent (Regression).



Features Scaling:

- **Min-Max Normalization:**

(Sometimes just called normalization)

- It scales each variable/feature in the $[0,1]$ range.
- This method preserves the shape of the original distribution and is sensitive to outliers.

```
from sklearn.preprocessing import MinMaxScaler
```

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

Features Scaling:

- **Mean Normalization:** (Sometimes just called standardization)
 - It produces a distribution centered at 0 with a standard deviation of 1.
 - This method “makes” a feature normally distributed. With outliers, the data will be scaled to a small interval.

`from sklearn.preprocessing import StandardScaler`

$$x' = \frac{x - \bar{x}}{\sigma}$$

Features Scaling:

- **Robust Scaling**

- All distributions have most of their densities around 0 and a shape that is more or less the same.
- The Interquartile range makes this method robust to outliers (hence the name).

```
from sklearn.preprocessing import RobustScaler
```

$$x' = \frac{x - Q_2(x)}{Q_3(x) - Q_1(x)}$$

, where Q are quartiles.

Batch/Vanilla GD

- **The main advantages:**

- We can use fixed learning rate during training without worrying about learning rate decay.
- It has straight trajectory towards the minimum and it is guaranteed to converge in theory to the global minimum if the loss function is convex and to a local minimum if the loss function is not convex.
- It has unbiased estimate of gradients. The more the examples, the lower the standard error.

- **The main disadvantages:**

- Even though we can use vectorized implementation, it may still be slow to go over all examples especially when we have large datasets.
- Each step of learning happens after going over all examples where some examples may be redundant and don't contribute much to the update.

Resources

- <https://www.coursera.org/learn/machine-learning>
- <https://machinelearningmastery.com/analytical-vs-numerical-solutions-in-machine-learning/>
- https://www.youtube.com/watch?v=e6kf6DDQVYA&ab_channel=TreeSoftMatterTheory
- https://en.wikipedia.org/wiki/Mathematical_optimization
- <https://builtin.com/data-science/gradient-descent>
- <https://towardsdatascience.com/k-means-clustering-algorithm-applications-evaluation-methods-and-drawbacks-aa03e644b48a>
- <https://math.stackexchange.com/questions/2202545/why-using-squared-distances-in-the-cost-function-linear-regression>
- <https://towardsdatascience.com/optimization-loss-function-under-the-hood-part-ii-d20a239cde11>
- <https://www.mathsisfun.com/gradient.html>
- <https://en.wikipedia.org/wiki/Derivative>
- <https://www.mathsisfun.com/calculus/derivatives-introduction.html>
- [https://math.libretexts.org/Bookshelves/Calculus/Map%3A_Calculus_Early_Transcendentals_\(Stewart\)/14%3A_Partial_Derivatives/14.01%3A_Functions_of_Several_Variables](https://math.libretexts.org/Bookshelves/Calculus/Map%3A_Calculus_Early_Transcendentals_(Stewart)/14%3A_Partial_Derivatives/14.01%3A_Functions_of_Several_Variables)
- <https://slideplayer.com/slide/4753135/>
- <https://en.wikipedia.org/wiki/Gradient>
- <https://www.khanacademy.org/math/multivariable-calculus/multivariable-derivatives/partial-derivative-and-gradient-articles/a/the-gradient>
- <https://la.mathworks.com/help/matlab/ref/meshc.html>
- <http://www.adeveloperdiary.com/data-science/how-to-visualize-gradient-descent-using-contour-plot-in-python/>
- https://rpubs.com/mgswiss15/M6C_7Multivariate
- <https://stats.stackexchange.com/questions/354046/coordinate-descent-with-constraints>
- <https://www.mathworks.com/help/optim/ug/local-vs-global-optima.html#:~:text=A%20local%20minimum%20of%20a,at%20all%20other%20feasible%20points>
- https://en.wikipedia.org/wiki/Maxima_and_minima
- <https://wngaw.github.io/linear-regression/>
- <http://www.cheerml.com/saddle-points>
- <https://towardsdatascience.com/understand-convexity-in-optimization-db87653bf920>
- <https://towardsdatascience.com/understand-convexity-in-optimization-db87653bf920>
- <https://www.sciencedirect.com/topics/engineering/convex-function>
- <https://www.math24.net/convex-functions#example2>
- <https://tutorial.math.lamar.edu/Classes/CalcI/NewtonsMethod.aspx>
- https://en.wikipedia.org/wiki/Newton's_method
- <https://tutorial.math.lamar.edu/Classes/CalcI/NewtonsMethod.aspx>

Resources

- <https://realpython.com/linear-regression-in-python/>
- <https://towardsdatascience.com/linear-regression-using-python-b136c91bf0a2>
- <https://towardsdatascience.com/why-norms-matters-machine-learning-3f08120af429>
- <https://towardsdatascience.com/why-norms-matters-machine-learning-3f08120af429>
- <https://machinelearningmastery.com/vector-norms-machine-learning/>
- <https://medium.com/linear-algebra/part-18-norms-30a8b3739bb>
- <https://heartbeat.fritz.ai/5-regression-loss-functions-all-machine-learners-should-know-4fb140e9d4b0>
- Andrew Ng, Machine Learning, Stanford University, Coursera
- <https://heartbeat.fritz.ai/5-regression-loss-functions-all-machine-learners-should-know-4fb140e9d4b0>
- <https://medium.com/data-science-365/linear-regression-with-gradient-descent-895bb7d18d52>
- https://www.holehouse.org/mlclass/17_Large_Scale_Machine_Learning.html
- <https://towardsdatascience.com/machine-learning-fundamentals-via-linear-regression-41a5d11f5220>
- <https://towardsdatascience.com/machine-learning-fundamentals-via-linear-regression-41a5d11f5220>
- <https://www.analyticsvidhya.com/blog/2019/08/detailed-guide-7-loss-functions-machine-learning-python-code/>
- <https://builtin.com/data-science/gradient-descent>
- <https://www.mltut.com/stochastic-gradient-descent-a-super-easy-complete-guide/>
- <https://towardsdatascience.com/linear-regression-using-gradient-descent-97a6c8700931>
- <https://kaigangi72.medium.com/stochastic-gradient-descent-demystified-part-1-8e4b897079b7>
- <https://medium.datadriveninvestor.com/gradient-descent-algorithm-b4c5afb4eb98>
- <https://medium.com/mindorks/an-introduction-to-gradient-descent-7b0c6d9e49f6>
- <https://medium.com/@venkatavinay222/at-the-end-machine-learning-is-all-about-optimization-ft-gradient-descent-e1588b7d95d2>
- “Hands-on Machine Learning with Scikit-Learn, Keras, and TensorFlow” by Aurélien Géron
- <https://laptrinhx.com/feature-scaling-why-and-how-3308094292/>
- <https://towardsdatascience.com/gradient-descent-algorithm-and-its-variants-10f652806a3>